

ANZ Analysis: Correlation between annual salary and age

In this analysis, we build a simple regression model and decision-tree model to predict the annual salary for each customer using the attribute age.

Data Cleansing

We need to remove duplicates and handle the missing values.

```
import pandas as pd
import numpy as np

data = pd.read_excel('~\Downloads\ANZ\ANZ synthesised transaction dataset.xlsx', delimiter='t')

print(data.head())
print(data.isnull().sum())

data.fillna(data.mean(), inplace=True)
data[['bpay_biller_code', 'merchant_id', 'merchant_suburb', 'merchant_state', 'merchant_long_lat']].fillna(0, inplace=True)

print(data.isnull().sum())

writer = pd.ExcelWriter('cleaned.xlsx')
data.to_excel(writer)
writer.save()print('Dataframe is written succesfully to excel file yo')
```

Code explanation:

1. `data.fillna(data.mean(), inplace=True)`: we fill the missing values using average of each integer/float columns.
2. `data[['bpay_biller_code', 'merchant_id', 'merchant_suburb', 'merchant_state', 'merchant_long_lat']].fillna(0, inplace=True)`: since these columns are string type, we can not find the average value of these. So we fill the missing with `0`.
3. This is optional, I choose to save the cleaned data into excel document.

Feature Engineering

In this part, we need to select from raw data into a form which is easier to interpret.

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt

data = pd.read_excel('~\Downloads\ANZ\cleaned.xlsx', index='customer_id')

data.drop_duplicates(['customer_id'], keep='first', inplace=True)
data['annual_salary'] = data['amount'] * 52

salary = data[data['txn_description'] == 'PAY/SALARY']
age_salary = salary[['age', 'annual_salary']]
```

Code explanation:

1. `data.drop_duplicates(['customer_id'], keep='first', inplace=True)`: we need to select distinct customer first.
2. `data['annual_salary'] = data['amount'] * 52`: in this data, customer received certain amount of salary by weekly, so we need to multiply with 52 weeks to get the annual salary.
3. `salary = data[data['txn_description'] == 'PAY/SALARY']`: create a subset `age_salary` which contains **PAY/SALARY** category of transactions and slice the column `age` and `annual_salary`.

Simple Regression Model

First, we build a simple regression model with the data above.

```
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error
from sklearn.metrics import mean_absolute_error

X = age_salary.iloc[:, :-1].values
y = age_salary.iloc[:, 1].values
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 1/5, random_state=0)

regressor = LinearRegression()
regressor.fit(X_train, y_train)

ax_train = plt
ax_train.scatter(X_train, y_train, color='red')
ax_train.plot(X_train, regressor.predict(X_train), color='blue')
ax_train.title('Annual Salary vs Customer Age (Training Set)')
ax_train.xlabel('Age')
ax_train.ylabel('Annual Salary')
ax_train.show()

ax_test = plt
ax_test.scatter(X_test, y_test, color='red')
ax_test.plot(X_test, regressor.predict(X_test), color='blue')
ax_test.title('Annual Salary vs Customer Age (Test Set)')
ax_test.xlabel('Age')
ax_test.ylabel('Annual Salary')
ax_test.show()

#y_pred = regressor.predict([[27]])
y_pred = regressor.predict(X_test)



#evaluate
lin_mse = mean_squared_error(y_pred, y_test)
lin_rmse = np.sqrt(lin_mse)
lin_mae = mean_absolute_error(y_pred, y_test)

print('Slope: ', regressor.coef_)
print('Intercept: ', regressor.intercept_)
print('Liner Regression R squared: %.4f' % regressor.score(X_test, y_test))
print('Liner Regression RMSE: %.4f' % lin_rmse)
print('Liner Regression MAE: %.4f' % lin_mae)
```

Code explanation:

1. `x`: the first column which contains Age array
`y`: the last column which contains Annual Salary array
2. `test_size=1/5`: we will split our dataset into 2 parts (data training and data test) and the ratio of **data test** compare to dataset is 1/5.
3. `regressor = LinearRegression()`: our training model will implement the Linear Regression.
`regressor.fit`: in this line, we pass the `X_train` which contains value of **Age** and `y_train` which contains values of **Annual Salary** to form up the model. This is the training process.
4. Now we can use it to predict *any values of X depends on y* or *any values of y depends on X*. First, we want to know Annual Salary for age 27: `y_pred = regressor.predict([[27]])`. Then we apply to all `X_test`.
5. The model need to evaluate, so we calculate *Root Mean Squared Error* and *Mean Absolute Error*. Here are the values:

```
Slope: [167.53325482]
Intercept: 97475.68457173448
Liner Regression R squared: -0.6984
Liner Regression RMSE: 38498.1724
Liner Regression MAE: 34563.7152
```



Decision Tree Model

Next we will build Decision Tree model.

```
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeRegressor
from sklearn.metrics import mean_squared_error

X = age_salary.iloc[:, :-1].values
y = age_salary.iloc[:, 1].values
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.20, random_state=0)

regressor = DecisionTreeRegressor()
regressor.fit(X_train, y_train)

#y_pred = regressor.predict([[27]])
y_pred = regressor.predict(X_test)

plt.scatter(X, y, color = 'red')
plt.plot(X, regressor.predict(X), color = 'blue')
plt.title('Annual Salary vs Customer Age')
plt.xlabel('Age')
plt.ylabel('Salary')
plt.show()

#evaluate
lin_mse = mean_squared_error(y_pred, y_test)
lin_rmse = np.sqrt(lin_mse)
print('Liner Regression RMSE: %.4f' % lin_rmse)
```

Code explanation:

1. `regressor = DecisionTreeRegressor()` `regressor.fit(X_train, y_train)` where we use decision tree model.
2. We also have the data test 20% `test_size = 0.20`
3. The RMSE for decision tree: `49587.8135`

